

EXPERIMENT 1

EXPERIMENT OBJECTIVE

To implement a Fully Connected Neural Network (FCNN) for classifying handwritten digits from the MNIST Dataset using NumPy.

DATA PREPROCESSING

Loading the MNIST Dataset

- The dataset is loaded from binary files containing images and labels.
- The images are 28x28 grayscale images, reshaped into a (784,) vector.
- The labels are converted into one-hot encoded vectors.

Data Augmentation

- **Random Rotation:** Images are rotated within a range of -15° to 15° with a 50% probability.
- **Horizontal Flip:** Images have a 50% chance of being flipped horizontally.

Splitting the Dataset

- The dataset is divided into training, validation, and test sets.
- The training set is further split into 80% training and 20% validation.

NEURAL NETWORK IMPLEMENTATION

Architecture

- **Input Layer:** 784 neurons (28x28 pixels flattened)
- **Hidden Layer 1:** 256 neurons, ReLU activation
- **Hidden Layer 2:** 128 neurons, ReLU activation
- **Output Layer:** 10 neurons (digits 0-9), softmax activation

Weight Initialization

- Weights are initialized using He (Kaiming) initialization.
- Biases are initialized to zeros.

Activation Functions

- **ReLU (Rectified Linear Unit):** Used in hidden layers.
- **Softmax:** Applied to the output layer for probability distribution.

Regularization

- **Dropout:** Randomly drops activations during training to prevent overfitting.
- **Gradient Clipping:** Limits gradient values to avoid exploding gradients.

TRAINING CONFIGURATION

Training the Model

- **Loss Function:** Cross-entropy loss is used.
- **Optimizer:** The model updates weights using backpropagation and gradient descent.
- **Learning Rate:** 0.2 (with decay over time)
- **Epochs:** Trained for 2500 epochs.
- **Batch Processing:** Mini-batch gradient descent is implemented.
- **Best Model Selection:** Saves weights of the best-performing model (lowest validation loss).

Model Checkpointing

- The best model weights (based on validation loss) are saved periodically to bestWeights.npy.

TRAINING AND VALIDATION RESULTS

Key Performance Metrics from Training Output

Epoch	Training Loss	Validation Loss	Accuracy (%)
0	1.9218	1.9253	30.28%
8	1.5442	1.5684	52.47%
31	1.1557	1.1522	66.05%
122	0.6481	0.6511	81.34%
249	0.4502	0.4500	87.50%
604	0.2744	0.2625	93.67%
1017	0.1331	0.1327	96.69%
1506	0.0542	0.0508	98.67%
2293	0.0135	0.0132	99.57%
2368	0.0086	0.0081	99.69%

Evaluation Results

- After training, the best model weights are loaded and tested on unseen test data.
- **Final Test Accuracy:** ~95.79%
- **Final Test Loss:** ~0.302

MODEL SAVING AND LOADING

- **Saving Weights:** The best model weights (lowest validation loss) are saved to disk.
- **Loading Weights:** Enables reloading the best weights for inference or further training.

RESULTS AND CONCLUSIONS

- The model achieves high accuracy using a simple fully connected architecture.
- Data augmentation and regularization significantly improve generalization.
- The saved best weights allow for consistent reproducibility of results.